

Import Data

```
import zipfile

# Download zip file
!wget https://storage.googleapis.com/ztm_tf_course/food_vision/pizza_steak.zip

# Unzip file
zip_ref = zipfile.ZipFile('pizza_steak.zip')
zip_ref.extractall()
zip_ref.close()

--2022-09-21 03:59:58-- https://storage.googleapis.com/ztm\_tf\_course/food\_vision/pizza\_steak.zip
Resolving storage.googleapis.com (storage.googleapis.com)... 108.177.127.128, 142.251.148.128
Connecting to storage.googleapis.com (storage.googleapis.com)|108.177.127.128|:443... connected
HTTP request sent, awaiting response... 200 OK
Length: 109540975 (104M) [application/zip]
Saving to: 'pizza_steak.zip'

pizza_steak.zip      100%[=====>] 104.47M  57.3MB/s   in 1.8s

2022-09-21 04:00:01 (57.3 MB/s) - 'pizza_steak.zip' saved [109540975/109540975]
```



Inspect the File

```
!ls pizza_steak
```

```
test  train
```

```
# Identify the path
```

```
!ls pizza_steak/train/steak/
```

1400700.jpg	1970100.jpg	2534307.jpg	3115772.jpg	3671077.jpg	703330.jpg
1403005.jpg	1984271.jpg	2535431.jpg	3116018.jpg	368073.jpg	703909.jpg
1404770.jpg	1987213.jpg	2535456.jpg	3128952.jpg	368162.jpg	704316.jpg
140832.jpg	1987639.jpg	2538000.jpg	3130412.jpg	368170.jpg	714298.jpg
141056.jpg	1995118.jpg	2543081.jpg	3136.jpg	3693649.jpg	720060.jpg
141135.jpg	1995252.jpg	2544643.jpg	313851.jpg	3700079.jpg	726083.jpg
1413972.jpg	199754.jpg	2547797.jpg	3140083.jpg	3704103.jpg	728020.jpg
1421393.jpg	2002400.jpg	2548974.jpg	3140147.jpg	3707493.jpg	732986.jpg
1428947.jpg	2011264.jpg	2549316.jpg	3142045.jpg	3716881.jpg	734445.jpg
1433912.jpg	2012996.jpg	2561199.jpg	3142618.jpg	3724677.jpg	735441.jpg
143490.jpg	2013535.jpg	2563233.jpg	3142674.jpg	3727036.jpg	740090.jpg
1445352.jpg	2017387.jpg	256592.jpg	3143192.jpg	3727491.jpg	745189.jpg
1446401.jpg	2018173.jpg	2568848.jpg	314359.jpg	3736065.jpg	752203.jpg
1453991.jpg	2020613.jpg	2573392.jpg	3157832.jpg	37384.jpg	75537.jpg
1456841.jpg	2032669.jpg	2592401.jpg	3159818.jpg	3743286.jpg	756655.jpg
146823.jpg	203450.jpg	2599817.jpg	3162376.jpg	3745515.jpg	762210.jpg



140833.jpg	203430.jpg	2599817.jpg	3102370.jpg	3743313.jpg	702210.jpg
1476404.jpg	2034628.jpg	2603058.jpg	3168620.jpg	3750472.jpg	763690.jpg
1485083.jpg	2036920.jpg	2606444.jpg	3171085.jpg	3752362.jpg	767442.jpg
1487113.jpg	2038418.jpg	2614189.jpg	317206.jpg	3766099.jpg	786409.jpg
148916.jpg	2042975.jpg	2614649.jpg	3173444.jpg	3770370.jpg	80215.jpg
149087.jpg	2045647.jpg	2615718.jpg	3180182.jpg	377190.jpg	802348.jpg
1493169.jpg	2050584.jpg	2619625.jpg	31881.jpg	3777020.jpg	804684.jpg
149682.jpg	2052542.jpg	2622140.jpg	3191589.jpg	3777482.jpg	812163.jpg
1508094.jpg	2056627.jpg	262321.jpg	3204977.jpg	3781152.jpg	813486.jpg
1512226.jpg	2062248.jpg	2625330.jpg	320658.jpg	3787809.jpg	819027.jpg
1512347.jpg	2081995.jpg	2628106.jpg	3209173.jpg	3788729.jpg	822550.jpg
1524526.jpg	2087958.jpg	2629750.jpg	3223400.jpg	3790962.jpg	823766.jpg
1530833.jpg	2088030.jpg	2643906.jpg	3223601.jpg	3792514.jpg	827764.jpg
1539499.jpg	2088195.jpg	2644457.jpg	3241894.jpg	379737.jpg	830007.jpg
1541672.jpg	2090493.jpg	2648423.jpg	3245533.jpg	3807440.jpg	838344.jpg
1548239.jpg	2090504.jpg	2651300.jpg	3245622.jpg	381162.jpg	853327.jpg
1550997.jpg	2125877.jpg	2653594.jpg	3247009.jpg	3812039.jpg	854150.jpg
1552530.jpg	2129685.jpg	2661577.jpg	3253588.jpg	3829392.jpg	864997.jpg
15580.jpg	2133717.jpg	2668916.jpg	3260624.jpg	3830872.jpg	885571.jpg
1559052.jpg	2136662.jpg	268444.jpg	326587.jpg	38442.jpg	907107.jpg
1563266.jpg	213765.jpg	2691461.jpg	32693.jpg	3855584.jpg	908261.jpg
1567554.jpg	2138335.jpg	2706403.jpg	3271253.jpg	3857508.jpg	910672.jpg
1575322.jpg	2140776.jpg	270687.jpg	3274423.jpg	386335.jpg	911803.jpg
1588879.jpg	214320.jpg	2707522.jpg	3280453.jpg	3867460.jpg	91432.jpg
1594719.jpg	2146963.jpg	2711806.jpg	3298495.jpg	3868959.jpg	914570.jpg
1595869.jpg	215222.jpg	2716993.jpg	330182.jpg	3869679.jpg	922752.jpg
1598345.jpg	2154126.jpg	2724554.jpg	3306627.jpg	388776.jpg	923772.jpg
1598885.jpg	2154779.jpg	2738227.jpg	3315727.jpg	3890465.jpg	926414.jpg
1600179.jpg	2159975.jpg	2748917.jpg	331860.jpg	3894222.jpg	931356.jpg
1600794.jpg	2163079.jpg	2760475.jpg	332232.jpg	3895825.jpg	937133.jpg
160552.jpg	217250.jpg	2761427.jpg	3322909.jpg	389739.jpg	945791.jpg
1606596.jpg	2172600.jpg	2765887.jpg	332557.jpg	3916407.jpg	947877.jpg
1615395.jpg	2173084.jpg	2768451.jpg	3326734.jpg	393349.jpg	952407.jpg
1618011.jpg	217996.jpg	2771149.jpg	3330642.jpg	393494.jpg	952437.jpg
1619357.jpg	2193684.jpg	2779040.jpg	3333128.jpg	398288.jpg	955466.jpg
1621763.jpg	220341.jpg	2788312.jpg	3333735.jpg	40094.jpg	9555.jpg
1623325.jpg	22080.jpg	2788759.jpg	3334973.jpg	401094.jpg	961341.jpg
1624450.jpg	2216146.jpg	2796102.jpg	3335013.jpg	401144.jpg	97656.jpg
1624747.jpg	2222018.jpg	280284.jpg	3335267.jpg	401651.jpg	979110.jpg
1628861.jpg	2223787.jpg	2807888.jpg	3346787.jpg	405173.jpg	980247.jpg
1632774.jpg	2230959.jpg	2815172.jpg	3364420.jpg	405794.jpg	982988.jpg
1636831.jpg	2232310.jpg	2818805.jpg	336637.jpg	40762.jpg	987732.jpg
1645470.jpg	2233395.jpg	2823872.jpg	3372616.jpg	413325.jpg	996684.jpg

```
import os
```

```
# list number of files in directory
```

```
for dirpath,dirnames,filenames in os.walk('pizza_steak'):
```

```
    print(f"There are {len(dirnames)} directories and {len(filenames)} in '{dirpath}'.")
```

```
    There are 2 directories and 0 in 'pizza_steak'.
```

```
    There are 2 directories and 0 in 'pizza_steak/test'.
```

```
    There are 0 directories and 250 in 'pizza_steak/test/steak'.
```

```
    There are 0 directories and 250 in 'pizza_steak/test/pizza'.
```

```
    There are 2 directories and 0 in 'pizza_steak/train'.
```

```
There are 0 directories and 750 in 'pizza_steak/train/steak'.
There are 0 directories and 750 in 'pizza_steak/train/pizza'.
```

```
num_steak_images_train = len(os.listdir('pizza_steak/train/steak'))

# Number of steak training images
num_steak_images_train

750

import pathlib
import numpy as np
data_dir = pathlib.Path('pizza_steak/train')
class_names = np.array(sorted([item.name for item in data_dir.glob('*')]))

print(class_names)

['pizza' 'steak']
```

View an image

```
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import random

def view_random_image(target_dir, target_class):
    # create target directory
    target_folder = target_dir + '/' + target_class

    random_image = random.sample(os.listdir(target_folder), 1)

    img = mpimg.imread(target_folder + '/' + random_image[0])

    plt.imshow(img)

    plt.title(target_class)
    plt.axis('off');

    print(f"Image shape: {img.shape}")
    return img

# view a random image from the data set
img = view_random_image(target_dir = 'pizza_steak/train/',
                        target_class='steak')
```

Image shape: (512, 512, 3)

steak




```
plt.figure()
plt.subplot(1,2,1)
steak_img = view_random_image('pizza_steak/train/','steak') # View an image from the steak tr
plt.subplot(1,2,2)
pizza_img = view_random_image('pizza_steak/train/','pizza') # View an image from the pizza tr
```

Image shape: (384, 512, 3)

Image shape: (512, 512, 3)

steak



pizza



```
# view image data
import tensorflow as tf
tf.constant(img)
```

```
<tf.Tensor: shape=(512, 512, 3), dtype=uint8, numpy=
array([[165, 117,  71],
       [172, 124,  78],
       [176, 129,  85],
       ...,
       [138,  14,   4],
       [154,  41,  23],
       [189,  84,  62]],

       [[159, 108,  65],
       [161, 110,  67],
       [159, 110,  67],
       ...,
       [148,  27,  16],
       [177,  70,  50],
```

```

[219, 119, 95]],

[[134, 78, 41],
 [132, 78, 40],
 [133, 79, 41],
 ...,
 [163, 51, 37],
 [207, 108, 85],
 [247, 157, 130]],

...,

[[ 16, 5, 3],
 [ 16, 5, 3],
 [ 16, 6, 4],
 ...,
 [ 5, 1, 0],
 [ 5, 1, 0],
 [ 6, 2, 1]],

[[ 14, 3, 1],
 [ 14, 3, 1],
 [ 15, 4, 2],
 ...,
 [ 6, 2, 1],
 [ 6, 2, 1],
 [ 6, 2, 1]],

[[ 11, 0, 0],
 [ 11, 0, 0],
 [ 12, 1, 0],
 ...,
 [ 6, 2, 1],
 [ 7, 3, 2],
 [ 7, 3, 2]]], dtype=uint8)>

```

```
img.shape
```

```
(512, 512, 3)
```

```
# Normalized image data
img/255.
```

```

array([[0.64705882, 0.45882353, 0.27843137],
       [0.6745098 , 0.48627451, 0.30588235],
       [0.69019608, 0.50588235, 0.33333333],
       ...,
       [0.54117647, 0.05490196, 0.01568627],
       [0.60392157, 0.16078431, 0.09019608],
       [0.74117647, 0.32941176, 0.24313725]],

       [[0.62352941, 0.42352941, 0.25490196],
       [0.63137255, 0.43137255, 0.2627451 ],
       [0.62352941, 0.43137255, 0.2627451 ],

```

```

...,
[0.58039216, 0.10588235, 0.0627451 ],
[0.69411765, 0.2745098 , 0.19607843],
[0.85882353, 0.46666667, 0.37254902]],

[[0.5254902 , 0.30588235, 0.16078431],
[0.51764706, 0.30588235, 0.15686275],
[0.52156863, 0.30980392, 0.16078431],
...,
[0.63921569, 0.2 , 0.14509804],
[0.81176471, 0.42352941, 0.33333333],
[0.96862745, 0.61568627, 0.50980392]],

...,

[[0.0627451 , 0.01960784, 0.01176471],
[0.0627451 , 0.01960784, 0.01176471],
[0.0627451 , 0.02352941, 0.01568627],
...,
[0.01960784, 0.00392157, 0. ],
[0.01960784, 0.00392157, 0. ],
[0.02352941, 0.00784314, 0.00392157]],

[[0.05490196, 0.01176471, 0.00392157],
[0.05490196, 0.01176471, 0.00392157],
[0.05882353, 0.01568627, 0.00784314],
...,
[0.02352941, 0.00784314, 0.00392157],
[0.02352941, 0.00784314, 0.00392157],
[0.02352941, 0.00784314, 0.00392157]],

[[0.04313725, 0. , 0. ],
[0.04313725, 0. , 0. ],
[0.04705882, 0.00392157, 0. ],
...,
[0.02352941, 0.00784314, 0.00392157],
[0.02745098, 0.01176471, 0.00784314],
[0.02745098, 0.01176471, 0.00784314]]])

```

Create train and test data sets

```

import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator

tf.random.set_seed(42)

# Preprocess data
train_datagen = ImageDataGenerator(rescale=1./255)
valid_datagen = ImageDataGenerator(rescale=1./255)

# Setup directories
train_dir = '/content/pizza_steak/train'

```

```
test_dir = 'pizza_steak/test'

train_data = train_datagen.flow_from_directory(directory=train_dir,
                                                batch_size=32,
                                                target_size=(224,224), # change image size/sha
                                                class_mode='binary',
                                                seed=42)

valid_data = valid_datagen.flow_from_directory(directory=test_dir,
                                                batch_size=32,
                                                target_size=(224,224),
                                                class_mode='binary',
                                                seed=42)
```

```
Found 1500 images belonging to 2 classes.
Found 500 images belonging to 2 classes.
```

Model 1 is a simple example of a neural network. This model performs no better than random chance.

```
tf.random.set_seed(42)

model_1 = tf.keras.Sequential([
    tf.keras.layers.Flatten(input_shape=(224,224,3)),
    tf.keras.layers.Dense(4, activation='relu'),
    tf.keras.layers.Dense(4, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])

#Compile model
model_1.compile(loss='binary_crossentropy',
                optimizer=tf.keras.optimizers.Adam(),
                metrics=['accuracy'])

# Fit model
history_1 = model_1.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=valid_data,
                        validation_steps=len(valid_data))

Epoch 1/5
47/47 [=====] - 13s 186ms/step - loss: 0.8251 - accuracy: 0.494
Epoch 2/5
47/47 [=====] - 9s 184ms/step - loss: 0.6932 - accuracy: 0.5006
Epoch 3/5
47/47 [=====] - 9s 184ms/step - loss: 0.6932 - accuracy: 0.5006
Epoch 4/5
47/47 [=====] - 9s 195ms/step - loss: 0.6932 - accuracy: 0.4827
```

Epoch 5/5
47/47 [=====] - 9s 183ms/step - loss: 0.6932 - accuracy: 0.5006



```
model_1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 150528)	0
dense (Dense)	(None, 4)	602116
dense_1 (Dense)	(None, 4)	20
dense_2 (Dense)	(None, 1)	5

=====
Total params: 602,141
Trainable params: 602,141
Non-trainable params: 0
=====

Model 2 increases the number of trainable parameters by increasing the dimensionality of the output space. This does significantly increase the accuracy of the model.

```
tf.random.set_seed(42)
```

```
model_2 = tf.keras.Sequential([  
    tf.keras.layers.Flatten(input_shape=(224,224,3)),  
    tf.keras.layers.Dense(100, activation='relu'),  
    tf.keras.layers.Dense(100, activation='relu'),  
    tf.keras.layers.Dense(100, activation='relu'),  
    tf.keras.layers.Dense(1, activation='sigmoid')  
])
```

```
model_2.compile(loss='binary_crossentropy',  
                optimizer=tf.keras.optimizers.Adam(),  
                metrics=['accuracy'])
```

```
history_2 = model_2.fit(train_data,  
                        epochs=5,  
                        steps_per_epoch=len(train_data),  
                        validation_data=valid_data,  
                        validation_steps=len(valid_data))
```

Epoch 1/5
47/47 [=====] - 10s 209ms/step - loss: 2.9554 - accuracy: 0.643
Epoch 2/5


```

47/47 [=====] - 8s 181ms/step - loss: 0.8232 - accuracy: 0.7307
Epoch 3/5
47/47 [=====] - 8s 180ms/step - loss: 0.9018 - accuracy: 0.7146
Epoch 4/5
47/47 [=====] - 9s 183ms/step - loss: 0.5541 - accuracy: 0.7687
Epoch 5/5
47/47 [=====] - 9s 182ms/step - loss: 0.4646 - accuracy: 0.7966

```



```
model_2.summary()
```

```
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
=====		
flatten_1 (Flatten)	(None, 150528)	0
dense_3 (Dense)	(None, 100)	15052900
dense_4 (Dense)	(None, 100)	10100
dense_5 (Dense)	(None, 100)	10100
dense_6 (Dense)	(None, 1)	101
=====		
Total params: 15,073,201		
Trainable params: 15,073,201		
Non-trainable params: 0		

```

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPool2D, Activation
from tensorflow.keras import Sequential
from pyparsing.core import StringEnd

```

Model 3 is a small 3 layer CNN

```

model_3 = Sequential([
    Conv2D(filters=10,
           kernel_size=3,
           strides=1,
           padding='valid',
           activation='relu',
           input_shape=(224,224,3)), # input layer
    Conv2D(10,3,activation='relu'),
    Conv2D(10,3,activation='relu'),
    Flatten(),
    Dense(1,activation='sigmoid') # output layer

```

```
])
```

```
model_3.compile(loss='binary_crossentropy',
                optimizer=Adam(),
                metrics=['accuracy'])
```

```
model_3.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 222, 222, 10)	280
conv2d_1 (Conv2D)	(None, 220, 220, 10)	910
conv2d_2 (Conv2D)	(None, 218, 218, 10)	910
flatten_2 (Flatten)	(None, 475240)	0
dense_7 (Dense)	(None, 1)	475241

=====
Total params: 477,341
Trainable params: 477,341
Non-trainable params: 0
=====


Model 3 improves on Model 2, even though it greatly reduces the number of trainable parameters.

```
history_3 = model_3.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=valid_data,
                        validation_steps=len(valid_data))
```

```
Epoch 1/5
47/47 [=====] - 19s 213ms/step - loss: 1.2145 - accuracy: 0.684
Epoch 2/5
47/47 [=====] - 9s 198ms/step - loss: 0.4326 - accuracy: 0.8146
Epoch 3/5
47/47 [=====] - 10s 202ms/step - loss: 0.3114 - accuracy: 0.879
Epoch 4/5
47/47 [=====] - 9s 198ms/step - loss: 0.1479 - accuracy: 0.9486
Epoch 5/5
47/47 [=====] - 9s 198ms/step - loss: 0.0694 - accuracy: 0.9847
```

```
import pandas as pd
pd.DataFrame(history_3.history).plot(figsize=(10,7)) #plot loss and accuracy of model
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f7ec84fbe50>



```
def plot_loss_curves(history):
    """
    Returns separate loss curves for training and validation metrics.
    """
    loss = history.history['loss']
    val_loss = history.history['val_loss']
    accuracy = history.history['accuracy']
    val_accuracy = history.history['val_accuracy']

    epochs = range(len(history.history['loss']))


    #plot loss
    plt.plot(epochs, loss, label='training_loss')
    plt.plot(epochs, val_loss, label='val_loss')
    plt.title('loss')
    plt.xlabel('epochs')
    plt.legend()

    #plot accuracy
    plt.figure()
    plt.plot(epochs, accuracy, label='training_accuracy')
    plt.plot(epochs, val_accuracy, label='val_accuracy')
    plt.title('accuracy')
```

```
plt.xlabel('epochs')
plt.legend();
```

```
plot_loss_curves(history_3)
```

#some overfitting appearing in this model



Model 4 adds in 2 MaxPool layers to reduce overfitting

```
model_4 = Sequential([
    Conv2D(10,3,activation='relu',input_shape=(224,224,3)),
    MaxPool2D(pool_size=2),
    Conv2D(10,3,activation='relu'),
    MaxPool2D(),
    Conv2D(10,3,activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(1,activation='sigmoid')
])
```

```
model_4.compile(loss='binary_crossentropy',
                optimizer=Adam(),
                metrics=['accuracy'])
```

```
history_4 = model_4.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=valid_data,
                        validation_steps=len(valid_data)
)
```

```
Epoch 1/5
47/47 [=====] - 10s 193ms/step - loss: 0.6293 - accuracy: 0.6293
Epoch 2/5
47/47 [=====] - 9s 188ms/step - loss: 0.4682 - accuracy: 0.7847
Epoch 3/5
47/47 [=====] - 9s 189ms/step - loss: 0.4247 - accuracy: 0.8126
Epoch 4/5
47/47 [=====] - 9s 189ms/step - loss: 0.4048 - accuracy: 0.8231
Epoch 5/5
47/47 [=====] - 9s 189ms/step - loss: 0.3892 - accuracy: 0.8306
```



```
model_4.summary()
```


Model: "sequential_3"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 222, 222, 10)	280
max_pooling2d (MaxPooling2D)	(None, 111, 111, 10)	0
conv2d_4 (Conv2D)	(None, 109, 109, 10)	910
max_pooling2d_1 (MaxPooling2D)	(None, 54, 54, 10)	0
conv2d_5 (Conv2D)	(None, 52, 52, 10)	910
max_pooling2d_2 (MaxPooling2D)	(None, 26, 26, 10)	0
flatten_3 (Flatten)	(None, 6760)	0
dense_8 (Dense)	(None, 1)	6761

```
Total params: 8,861
Trainable params: 8,861
Non-trainable params: 0
```

```
plot_loss_curves(history_4)
```

```
#this model slightly improved accuracy while reducing overfitting.
```



Data Augmentation: we will attempt to reduce overfitting even more by augmenting the training data

[illegible]

```

height_shift_range=0.2,
horizontal_flip=True)

train_datagen = ImageDataGenerator(rescale=1/255.)

test_datagen = ImageDataGenerator(rescale=1/255.)

print('Augmented training data')
train_data_augmented = train_datagen.flow_from_directory(train_dir,
                                                         target_size=(224,224),
                                                         batch_size=32,
                                                         class_mode='binary',
                                                         shuffle=False) #False for image example below,
print('Non-augmented training data')
train_data = train_datagen.flow_from_directory(train_dir,
                                              target_size=(224,224),
                                              batch_size=32,
                                              class_mode='binary',
                                              shuffle=False) #False for image example below,

print('Non-augmented test data')
test_data = test_datagen.flow_from_directory(test_dir,
                                             target_size=(224,224),
                                             batch_size=32,
                                             class_mode='binary')

Augmented training data
Found 1500 images belonging to 2 classes.
Non-augmented training data
Found 1500 images belonging to 2 classes.
Non-augmented test data
Found 500 images belonging to 2 classes.

images, labels = train_data.next()
augmented_images, augmented_labels = train_data_augmented.next()

#Randomly choose an image and show the original and augmented version
random_number = random.randint(0,31)
print(f'Showing image number: {random_number}')
plt.imshow(images[random_number])
plt.title(f'Original image')
plt.axis(False)
plt.figure()
plt.imshow(augmented_images[random_number])
plt.title(f'Augmented images')
plt.axis(False);

```

Showing image number: 14

Original image



Augmented images



```
#Reshuffling training data
```

```
train_data_augmented = train_datagen_augmented.flow_from_directory(train_dir,  
                                                                    target_size=(224,224),  
                                                                    batch_size=32,  
                                                                    class_mode='binary',  
                                                                    shuffle=True)  
  
train_data = train_datagen.flow_from_directory(train_dir,  
                                                                    target_size=(224,224),  
                                                                    batch_size=32,  
                                                                    class_mode='binary',  
                                                                    shuffle=True)
```

```
Found 1500 images belonging to 2 classes.
```

```
Found 1500 images belonging to 2 classes.
```

Model 5 is identical to Model 4 but uses augmented data

```
model_5 = Sequential([  
    Conv2D(10,3,activation='relu'),  
    MaxPool2D(pool_size=2),  
    Conv2D(10,3,activation='relu'),  
    MaxPool2D(pool_size=2),  
    ...  
])
```



```
        Conv2D(10,3,activation='relu'),
        MaxPool2D(pool_size=2),
        Flatten(),
        Dense(1,activation='sigmoid')
    ])
```

```
model_5.compile(loss='binary_crossentropy',
                optimizer=Adam(),
                metrics=['accuracy'])
```

```
history_5 = model_5.fit(train_data_augmented,
                        epochs=5,
                        steps_per_epoch=len(train_data_augmented),
                        validation_data=test_data,
                        validation_steps=len(test_data))
```

Epoch 1/5

47/47 [=====] - 22s 466ms/step - loss: 0.6557 - accuracy: 0.626

Epoch 2/5

47/47 [=====] - 22s 461ms/step - loss: 0.5181 - accuracy: 0.742

Epoch 3/5

47/47 [=====] - 22s 459ms/step - loss: 0.4895 - accuracy: 0.774

Epoch 4/5

47/47 [=====] - 22s 459ms/step - loss: 0.4744 - accuracy: 0.778


Epoch 5/5

47/47 [=====] - 22s 460ms/step - loss: 0.4516 - accuracy: 0.796



```
plot_loss_curves(history_5)
```





Model 6 no longer augments data and adds additional Conv2D layers. This structure matches the Tiny VGG: <https://poloclub.github.io/cnn-explainer/>

```
model_6 = Sequential([
    Conv2D(10, 3, activation='relu', input_shape=(224, 224, 3)), # same input shape as our image
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Conv2D(10, 3, activation='relu'),
    Conv2D(10, 3, activation='relu'),
    MaxPool2D(),
    Flatten(),
    Dense(1, activation='sigmoid')
])

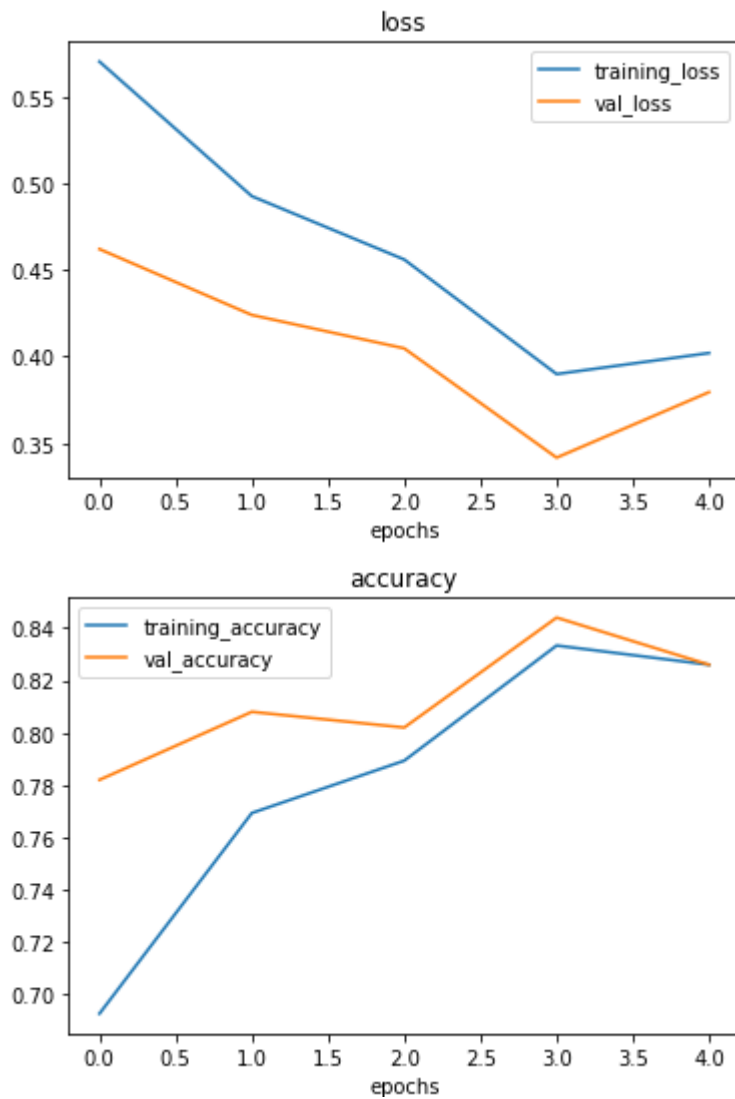
model_6.compile(loss="binary_crossentropy",
                optimizer=tf.keras.optimizers.Adam(),
                metrics=["accuracy"])

history_6 = model_6.fit(train_data,
                        epochs=5,
                        steps_per_epoch=len(train_data),
                        validation_data=test_data,
                        validation_steps=len(test_data))
```

```
Epoch 1/5
47/47 [=====] - 10s 199ms/step - loss: 0.5703 - accuracy: 0.693
Epoch 2/5
47/47 [=====] - 9s 194ms/step - loss: 0.4926 - accuracy: 0.7693
Epoch 3/5
47/47 [=====] - 9s 194ms/step - loss: 0.4560 - accuracy: 0.7893
Epoch 4/5
47/47 [=====] - 9s 196ms/step - loss: 0.3899 - accuracy: 0.8333
Epoch 5/5
47/47 [=====] - 9s 198ms/step - loss: 0.4020 - accuracy: 0.8266
```

```
plot_loss_curves(history_6)
```

```
# Overfitting is greatly reduced
```



model_6.summary()

Model: "sequential_5"

Layer (type)	Output Shape	Param #
conv2d_9 (Conv2D)	(None, 222, 222, 10)	280
conv2d_10 (Conv2D)	(None, 220, 220, 10)	910
max_pooling2d_6 (MaxPooling 2D)	(None, 110, 110, 10)	0
conv2d_11 (Conv2D)	(None, 108, 108, 10)	910
conv2d_12 (Conv2D)	(None, 106, 106, 10)	910
max_pooling2d_7 (MaxPooling 2D)	(None, 53, 53, 10)	0
flatten_5 (Flatten)	(None, 28090)	0
dense_10 (Dense)	(None, 1)	28091

=====
Total params: 31,101
Trainable params: 31,101
Non-trainable params: 0
=====

[Colab paid products](#) - [Cancel contracts here](#)

✓ 0s completed at 11:06 PM

